

SANYAM KAUL

kaulsanyam@gmail.com — +91-7018498338 — [LinkedIn](#) — [GitHub](#) — [LeetCode](#)

TECHNICAL SKILLS

- **Languages:** Java, Python, SQL, C++, Golang, Rust, Javascript, Shell/Bash.
- **AI/ML:** LLMs (RAG, CPT, SFT, RLHF, DPO, RLVR, LoRA/QLoRA), Transformers, LangChain, LangGraph, MCP.
- **Backend:** Spring Boot, FastAPI, REST, gRPC, GraphQL, Node.js, Microservices.
- **Data & Infrastructure:** Cloud (AWS, Azure), EKS, Lambda, SQS, SNS, S3, Docker, Kubernetes
- **Databases:** PostgreSQL, MongoDB, ChromaDB, Elasticsearch.
- **Tools:** Git, Jenkins, GitHub Actions, CI/CD, Grafana, Datadog, Prometheus.

EXPERIENCE

Netradyne

Data Engineer (ML Systems)

Bengaluru, India

Jul 2025 - Present

- Architected AI-powered agent for automating Root Cause Analysis workflow Recursive Language Model Architecture, and custom MCP servers, reducing analysis time by **60%**.
- Spearheaded the implementation of cloud monitoring System (Python, MongoDB, FastAPI) for machine metrics and costs, resulting in a **10–15% reduction in cloud spend**.
- Scaled AI Evals System using EKS, SQS, SNS, and ASGs, reducing latency for 50K+ video dataset creation and reprocessing, reducing total processin time by 35%.
- **Tech Stack:** Python, FastMCP, LangChain, LangGraph, FastAPI, MongoDB, PostgreSQL, Redis, AWS (EKS, S3, SQS, SNS, CloudWatch), Dash.

WinesecretHK — AI-powered Food and Beverage Sommelier

Remote

AI Architect

Apr 2024 - Jun 2025

- Designed, developed, deployed and maintained a multimodal **RAG-based AI Pairing System** using Python, AWS RDS, and ChromaDB for high-precision flavour-profile matching beyond standard semantic search.
- Implemented a multi-modal agent using RLM and cross-modal reasoning, achieving **70% more accurate pairings** than baseline semantic search and 20% improvement over standard hybrid RAG.
- Integrated Stripe payment systems and POS webhooks to support seamless ordering and transaction workflows.
- **Tech Stack:** Python, FastAPI, AWS RDS, ChromaDB, LangChain, LangGraph, Stripe.

BlackRock

Software Development Engineer 1

Gurugram, India

Jan 2021 - Jun 2023

- Development of a full-stack Financial Information Reporting Tool using Spring Boot and Angular, reducing client report setup time from **1.5 months to 2-3 days**.
- Standardized REST and GraphQL API workflows across **50+ global teams**, significantly improving system reliability and developer productivity.
- Built an **Azure DevOps** analytics dashboard to track CI/CD performance and improved backend test coverage to **85%+**.
- **Tech Stack:** Java (Spring Boot), Angular, REST, GraphQL, Golang, **Azure DevOps**.

PROJECTS

Car Insurance OCR (Structured Data Extraction) —

POC

- Developed a Python-based pipeline to extract structured data from unstructured insurance PDFs using Tesseract OCR and OpenAI LLMs.
- **Tech Stack:** Python, OpenAI APIs, Tesseract OCR, FastAPI.

Fraud Detection using Representation Learning —

GitHub

- Built a Variational Autoencoder (VAE) model for credit card fraud detection on imbalanced datasets, achieving **91%+ accuracy**.
- **Tech Stack:** Python, PyTorch, VAE, Scikit-Learn.

LooksMaxing AI Platform —

Live Link

- Developed an AI-driven personalization platform using Gemini and prompt engineering for virtual try-on and recommendation workflows.
- **Tech Stack:** Python, FastAPI, Gemini API, Prompt Engineering.

EDUCATION

IIT Hyderabad

M.Tech in Computer Science

Hyderabad, India

2025

- CGPA: 9 — Specialized in Distributed Systems, Blockchain and Applied AI.

VIT-AP University

B.Tech in Computer Science (Networks & Security)

Amaravati, India

2021

- CGPA: 9 — Focus on Cybersecurity and Network Infrastructure.

RESEARCH & ACHIEVEMENTS

- **M.Tech Thesis:** Optimized HyperLedger Fabric blockchain architecture to reduce node failure rates; presented at **Indosys 2025**.
[arXiv](#)
- **1st Place (IITH FoML):** Developed a winning ensemble model (Random Forest/GBDT) for predictive analytics, achieving 98% accuracy in predicting the correct user class on the test set.
- **15th/100+ (IITH DL):** Ranked in the top 15 in Deep Learning competition using ResNet and ViT architectures, achieving 92% accuracy in labeling image data on the test set.
- **Amazon AI Hackathon:** Developed a Vision-Language Model (VLM) based system to extract product data from images, featuring a chat-based UI for users to query product information.
- **Deep Learning Architectures:** Implemented CNN, RNN, and Transformers from scratch using NumPy to deepen core algorithmic understanding.